

Oracle SQL

By Dhandapani Yedappalli Krishnamurthi Sep 20, 2025

Part 1: Database & DBMS Fundamentals

1. Introduction to Database

Definition:

- A **database** is an organized collection of data, stored and retrieved electronically.
- Example: Banking database → tables for customers, accounts, transactions.

Illustration:

 Database = Digital cupboard →

- **Table** = Shelf
 - **Row** = File
 - **Column** = Label
-

2. Introduction to DBMS

DBMS (Database Management System):

- A software that helps store, retrieve, and manage data.
- Examples: Oracle, MySQL, SQL Server, PostgreSQL.

Advantages:

- ✓ Data consistency
- ✓ Security

- ✓ Backup & Recovery
 - ✓ Multi-user support
-

3. Characteristics of DBMS

- **Data Abstraction** (Users don't need to know low-level storage)
 - **Data Independence** (Schema changes don't affect apps)
 - **Concurrent Access** (Multiple users safely access data)
 - **Security** (User access control)
 - **Integrity Constraints** (Rules to ensure correct data)
-

4. DBMS Models

- **Hierarchical Model** – Tree structure (Parent → Child)
- **Network Model** – Records linked by relationships
- **Relational Model** – Tables with rows & columns (most common)

 Oracle uses Relational Model (RDBMS).

5. Relational DBMS (RDBMS)

- Stores data in **tables (relations)**.
- Each **row (tuple)** = a record.
- Each **column (attribute)** = a property.

Example in Oracle SQL:

```
CREATE TABLE Customers (  
    Cust_ID    NUMBER PRIMARY KEY,  
    Cust_Name  VARCHAR2(50),  
    City       VARCHAR2(30)  
);
```

Insert Data:

```
INSERT INTO Customers VALUES (1, 'Anita',  
'Chennai');  
INSERT INTO Customers VALUES (2, 'Rahul',  
'Bangalore');
```

Query Data:

```
SELECT * FROM Customers;
```

Cust_ID	Cust_Name	City
1	Anita	Chennai
2	Rahul	Bangalore

6. Data Integrity

Integrity ensures data correctness & consistency.

Types:

- **Entity Integrity:** Primary Key must be unique & not NULL.
- **Referential Integrity:** Foreign Key must reference existing record.
- **Domain Integrity:** Values must be valid for column datatype.

Example (Oracle SQL):

```
CREATE TABLE Orders (  
    Order_ID  NUMBER PRIMARY KEY,  
    Cust_ID   NUMBER REFERENCES Customers(Cust_ID),  
    Amount    NUMBER CHECK (Amount > 0)  
);
```

👉 Here:

- `Order_ID` ensures **entity integrity**.
 - `Cust_ID` ensures **referential integrity**.
 - `CHECK (Amount > 0)` ensures **domain integrity**.
-

7. Security in Database

- **Authentication** (Login with username/password)
- **Authorization** (Grant/Revoke rights)
- **Encryption** (Storing sensitive data securely)

Example (Oracle SQL):

```
CREATE USER trainee IDENTIFIED BY pass123;
GRANT CONNECT, RESOURCE TO trainee;
REVOKE RESOURCE FROM trainee;
```

8. Normalization & Codd's Rules

- **Normalization**: Process of organizing data to avoid redundancy.
- **Codd's Rules**: 12 golden rules for a **"fully relational" DBMS** (e.g., data should be stored in tables, not files).

👉 Example:

- **Unnormalized Data:**

Cust_ID	Cust_Name	Orders
1	Anita	0101, 0102

- **Normalized Data:**

Customers:	Cust_ID	Cust_Name
Orders:	Order_ID	Cust_ID

9. First Normal Form (1NF)

- No **repeating groups**.
- Each column must hold **atomic values**.

Example:

```
-- ❌ Bad (multi-valued Orders)
Cust_ID | Cust_Name | Orders
1       | Anita    | 0101, 0102

-- ✅ Good (split into atomic rows)
Cust_ID | Cust_Name | Order_ID
1       | Anita    | 0101
1       | Anita    | 0102
```

10. Second Normal Form (2NF)

- Table should be in 1NF.
- No **partial dependency** (non-key attribute depending on part of composite key).

11. Third Normal Form (3NF)

- Table in 2NF.
 - No **transitive dependency** (non-key depending on another non-key).
-

How SQL Works Under the Hood

When you run a SQL query like:

```
SELECT Cust_Name
FROM Customers
WHERE City = 'Chennai';
```

Oracle (or any RDBMS) goes through these steps:

1. SQL Parsing

- SQL is **parsed** by the SQL engine.
- Checks
 - **syntax** (**SELECT** written correctly?),
 - **semantics** (does **Customers** table exist? does **Cust_Name** column exist?).
- If invalid → you get an **error** immediately.

Oracle stores parsed SQL in the Shared Pool (library cache).

- If the same query comes again → it reuses the **execution plan** (saves time).

👉 Example:

```
SELECT * FROM Customers;  -- parsed and cached
```

If you run it again, Oracle won't re-parse.

2. Query Optimization

- Oracle checks **multiple ways** to run the query.
- Chooses the **best execution plan** using the **Cost-Based Optimizer (CBO)**.
- It considers:
 - Indexes (use them or not?)
 - Table size
 - Joins order
 - Statistics

Example:

```
EXPLAIN PLAN FOR
SELECT Cust_Name FROM Customers WHERE City = 'Chennai';

SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

This shows **how Oracle plans to run your query** (Index Scan, Full Table Scan, etc.).

3. Row Source Generation (Execution Plan)

- Oracle **creates a plan** step by step:
 - **Full Table Scan** (check all rows)
 - or **Index Scan** (jump directly to matching rows).

Example: If the **City** column has an index, Oracle may use **Index Range Scan** instead of scanning the whole table.

4. Query Execution

- Oracle actually **runs the chosen plan**.
- Retrieves data blocks from **Buffer Cache** (memory).
- If not in cache → fetch from **disk (datafiles)**.

👉 Oracle works in units called **blocks** (usually 8 KB).

5. Fetching Results

- Results are fetched row by row.
- Oracle sends results to the client in **batches (arrays)**, not one row at a time (to optimize performance).

Example: SQL*Plus might fetch 15 rows per batch by default.

Behind the Scenes Components

- **Parser** → breaks SQL into tokens.
 - **Optimizer** → decides best plan.
 - **Row Source Generator** → prepares step-by-step plan.
 - **Query Executor** → retrieves & returns rows.
 - **Buffer Cache** → memory area to reduce disk reads.
-

Example Walkthrough (with Index)

```
-- Assume index on City  
SELECT Cust_Name FROM Customers WHERE City = 'Chennai';
```

Under the hood:

1. Parser validates SQL.
 2. Optimizer checks → "Should I scan whole Customers table or use index on City?"
 3. Finds that **index on City** is efficient.
 4. Execution plan: **Index Range Scan → Table Access by RowID**.
 5. Fetches matching rows from buffer cache / disk.
 6. Sends results to user.
-

Analogy

Think of it like a **restaurant order system** :

- You (client) → give the **order** (SQL query).
 - Waiter (parser) → checks if order is valid.
 - Chef (optimizer) → decides the best way to cook it.
 - Kitchen (executor) → prepares the dish.
 - Waiter brings results (rows) to you.
-

Performance Tips

- **Indexes** help optimizer choose faster paths.
 - **Statistics** (ANALYZE TABLE or DBMS_STATS) help optimizer know table sizes.
 - **Bind variables** (:param) let Oracle reuse execution plans.
 - **Avoid SELECT *** → fetch only needed columns.
 - Use **EXPLAIN PLAN** and **AUTOTRACE** to see what's happening.
-

✓ In short:

SQL → Parse → Optimize → Generate Plan → Execute → Fetch results.